CS 4641: Supervised Learning

Bradley Reardon

February 10, 2019

1 Introduction

The purpose of this assignment is to evaluate various supervised learning techniques in the context of two classification algorithms. We will focus on analyzing and adjusting five learning algorithms using two publically-available data sets, through cross-validation and hyperparameter adjustments. Then, we can use our adjusted models to compare supervised learning techniques in the context of each of our classification algorithms.

2 Classification algorithms

For the purposes of this report, data sets were obtained from the UCI Machine Learning Repository. Each data set downloaded was processed with a custom script in the corresponding folder, **process.py**, which randomly separated the data into an 80% training data and 20% test data split using the **train_test_split** function from scikit-learn. The split data was then serialized using Python's **pickle** module into **.dataset** files, to ensure that the training and test sets remained constant for evaluation.

For the purposes of this assignment, cross-validation scoring of the various algorithms for each classification problem will be shown as the average score of 5-fold cross validation.

2.1 Car data set

The Car Evaluation Database was created in June, 1997 by Marko Bohanec. It contains 1728 instances and six attributes. The purpose of this database is to attempt to classify a car as an unacceptable, acceptable, good, or very good choice based on factors including cost of ownership, comfort, and safety. Full details about the data set can be found at the source link below.

Note that this specific data set contains only categorical attributes. As scikit-learn does not support non-continuous variables, a one-hot encoder was used to re-encode categorical features into multiple dimensions.

The problem at hand for this dataset is determining the acceptability for a car purchase based on the aforementioned attributes. Because this dataset notes that the instances in the data set completely cover the attribute space, this data set is interesting in particular due to its usefulness in comparing the optimization of different supervised learning algorithms.

Source: https://archive.ics.uci.edu/ml/datasets/car+evaluation

2.2   Breast Cancer Wisconsin data set

The Breast Cancer Wisconsin data set was donated to the UCI Machine Learning Repository in 1992, and contains data from one doctor's clinical cases, gathered from January 1989 to November 1991. In total, there are 699 instances signifying information about breast tumors such as clump thickness, uniformity in shape and size, and other screening factors. Data points are identified by two classes – benign or malignant. The features of the data points are encoded as 9 continuous attributes rating the screening factor from 1 to 10.

This data set contains unknowns in the form of question marks in the data. To deal with this, missing values were imputed, calculating missing data points based on the mean of other points in the specific column of the missing attribute.

The problem at hand for this dataset is determining whether a tumor is benign or malignant based on tumor screening characteristics identified in the data set.

Source: https://archive.ics.uci.edu/ml/datasets/Breast+Cancer+Wisconsin+%28Original%29

3   Decision trees with pruning

3.1   Parameter selection

To begin evaluating the decision tree algorithm, hyperparameter selection was essential to tuning the algorithms to each data set. Both the car and cancer datasets were evaluated with the GINI index criterion as well as the entropy criterion. The two criterion had no discernible differences in training results, so as a result, GINI was selected for the decision tree classifier on both sets, as it allows us to avoid a logarithmic calculation, saving compute time.

The next parameters that needed tuning were the pruning parameters included in the scikit-learn decision tree classifier. Namely, the max depth of the tree, as well as the minimum samples per leaf, were tuned to fit each data set.

For the car data set, parameter selection showed that a lower-depth tree had much lower accuracy, likely due to the fact that the dataset had instances covering the entire attribute space. As a result, a max depth of 9 nodes was selected for the car data set, with a minimum number of 5 samples occurring at each leaf node to minimize the amount of information loss.

The cancer data set, on the other hand, responded much better to more aggressive pruning. The pruned decision tree classifier responded to the data with the best F1 score with a max depth of 5 nodes, and a minimum number of 10 samples occurring at each leaf.

3.2   Performance

Performance of the tuned classifier for each data set was measured with both wall clock time for fitting the model with the training data, as well as making a prediction with the test data. For the breast cancer data set, pruning resulted in both higher precision and accuracy, with a performance gain during prediction over a non-pruned decision tree (from 0.801ms to 0.479ms). However, the

classifier took 0.57 seconds longer to train with pruning enabled (no pruning: 2.38s, pruning: 2.95s). For this data set, with the determined parameters was successful.

However, the car data set showed different results with pruning enabled. There was no significant change in either the precision or accuracy measures. The car data set's decision tree classifier took 11.51 seconds to fit without pruning, however with pruning this was reduced to 7.35s with pruning. Prediction runtime for the pruned tree was only slightly longer than the unpruned tree. Due to the completeness of the data, pruning is likely not necessary for this dataset without further tuning.

Learning curves for the data sets, with and without pruning enabled, can be found in Figure 1.

(a) Car data set, no pruning

(b) Car data set, with pruning

(c) Breast cancer data set, no pruning

(d) Breast cancer data set, with pruning
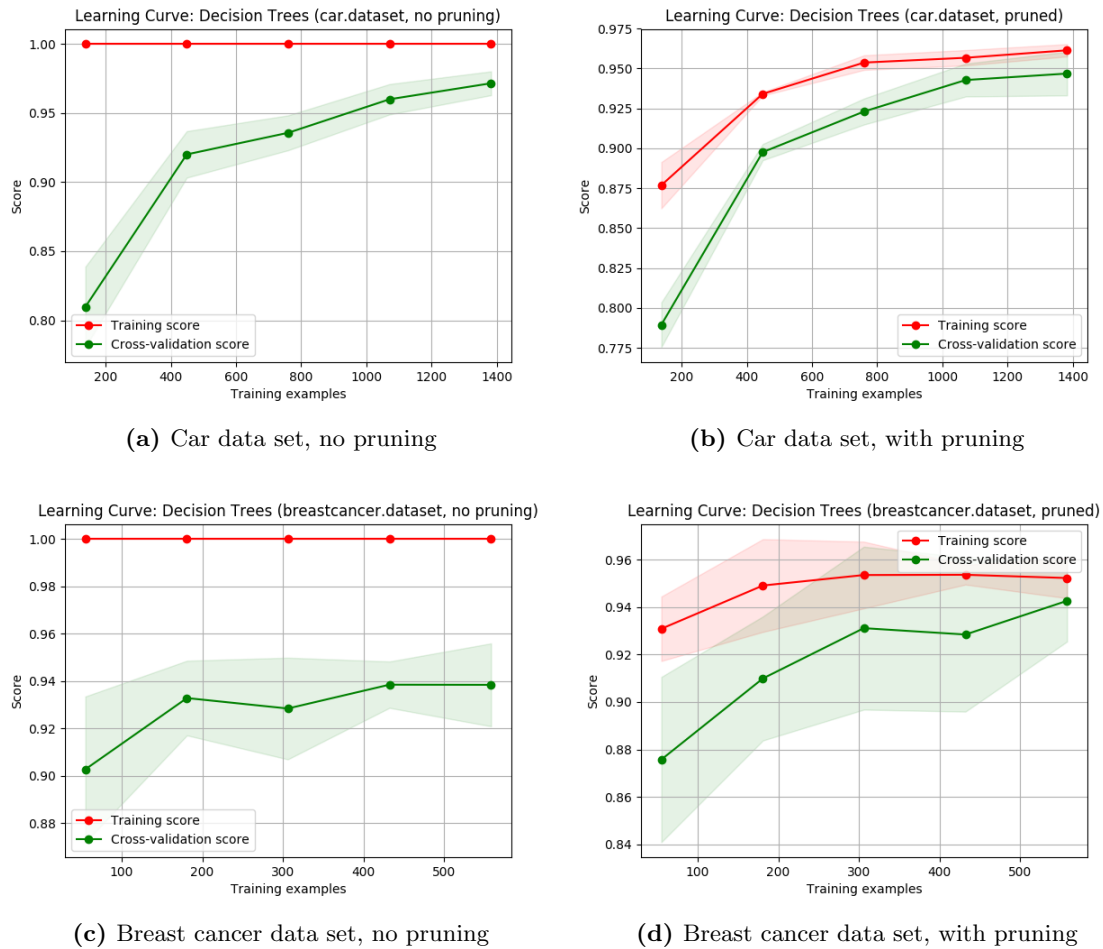
Figure 1: Learning curves for the car and breast cancer data sets using a decision tree classifier, with and without pruning.

4 Neural networks

TODO

5 Boosting

5.1 Parameter selection

TODO determining learning rate (l_rate) Figure 2.

car rate 1.0 n_est 100
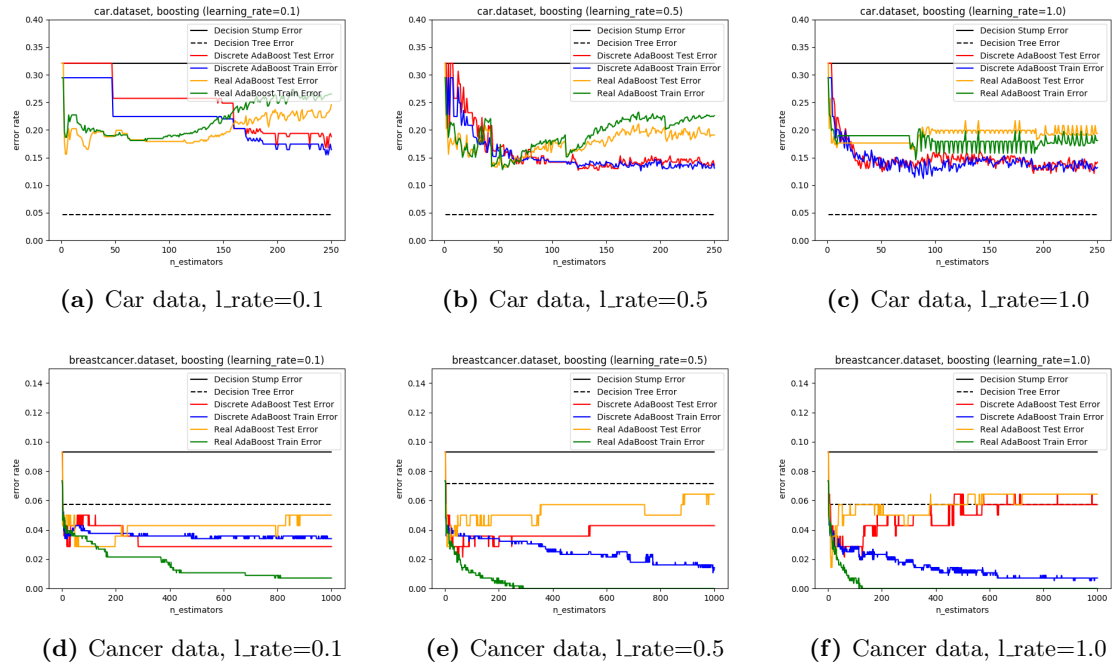
cancer rate 1.0 n_est 200

(a) Car data, l_rate=0.1   **(b)** Car data, l_rate=0.5   **(c)** Car data, l_rate=1.0

(d) Cancer data, l_rate=0.1   **(e)** Cancer data, l_rate=0.5   **(f)** Cancer data, l_rate=1.0

Figure 2: Selecting a learning rate and number of estimators for the AdaBoot algorithm, using the car and breast cancer sets.

6 Support vector machines

6.1 Performance

Learning curves for the car and breast cancer data sets can be found in Figure 3 on the following page.

TODO kernel params?

7 k-nearest neighbors

TODO

**(a)** Car data, linear kernel   **(b)** Car data, RBF kernel   **(c)** Car data, poly kernel

(d) Breast cancer, linear kernel   **(e)** Breast cancer, RBF kernel   **(f)** Breast cancer, poly kernel

Figure 3: Learning curves for the car and breast cancer data sets using an SVM classifier, with linear, RBF, and poly kernels.

7.1 Parameter selection

car k=11 cancer k=5

Lowest k with highest cross-validation

TODO Figure 4 on the next page

7.2 Performance

TODO Figure 5 on the following page

8 Analysis

TODO

9 Discussion

TODO

(a) Car data set

(b) Breast cancer data set

Figure 4: Testing k-nearest neighbors with various values of k.

(a) Car data set, $k = 11$

(b) Breast cancer data set, $k = 5$

Figure 5: Learning curves for k-nearest neighbors with optimized k-values.

10   Conclusion

TODO